

✓ Final Machine Learning Project

I downloaded data from a Kaggle competition titled "**House Prices: Advanced Regression Techniques**".

The goal is to build a model that predicts house prices based on their characteristics.

The project will be built by testing several models, selecting the best one, and then creating a simple web application to display the predictions directly.

Data Download Link

<https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques/data>

✓ [1] Import Required Libraries

```
1 #####
2 # 1. Suppress Warnings #
3 #####
4 import warnings
5 warnings.filterwarnings('ignore')
6
7
8 #####
9 # 2. Core Libraries #
10 #####
11 import os
12 import joblib
13 import numpy as np
14 import pandas as pd
15
16
17 #####
18 # 3. Data Visualization #
19 #####
20 import seaborn as sns
21 import matplotlib.pyplot as plt
22
23 %matplotlib inline
24 sns.set_palette("Set2")
25 sns.set_style("whitegrid")
26
27 pd.set_option('display.width', 1000)
28 pd.set_option('display.max_columns', None)
29 pd.set_option('display.max_colwidth', None)
30
31
32 #####
33 # 4. Data Preprocessing #
34 #####
35 from sklearn.impute import SimpleImputer
36 from sklearn.compose import ColumnTransformer
37 from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder, PolynomialFeatures
38
39
40 #####
41 # 5. Machine Learning Models #
42 #####
43 from sklearn.svm import SVR
44 from sklearn.dummy import DummyRegressor
45 from sklearn.tree import DecisionTreeRegressor
46 from sklearn.neural_network import MLPRegressor
47 from sklearn.neighbors import KNeighborsRegressor
48 from sklearn.linear_model import LinearRegression, Ridge, Lasso
49 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
50
51 # Advanced models
52 from xgboost import XGBRegressor
53 from lightgbm import LGBMRegressor
54
55
56 #####
57 # 6. Pipeline & Data Splitting #
58 #####
```

```

59 from sklearn.pipeline import Pipeline
60 from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
61
62
63 ##### # 7. Random Distributions (for Randomized Search) #
64 ##### # 8. Evaluation Metrics #
65 #####
66 from scipy.stats import randint, uniform
67
68
69 #####
70 # 8. Evaluation Metrics #
71 #####
72 from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
73
74
75 #####
76 # 9. Additional Tools #
77 #####
78 from google.colab import files
79
80
81 #####
82 # 10. Deployment #
83 #####
84 !pip install gradio --quiet
85 import gradio as gr
86
87 print("All required libraries have been successfully loaded!")

```

→ All required libraries have been successfully loaded!

▼ [2] Load & Read Dataset Files

```
1 uploaded = files.upload()
```

→ Choose Files | 2 files

- test.csv(text/csv) - 451405 bytes, last modified: 12/15/2019 - 100% done
- train.csv(text/csv) - 460676 bytes, last modified: 12/15/2019 - 100% done

Saving test.csv to test (2).csv
Saving train.csv to train (2).csv

```
1 df_train = pd.read_csv('train.csv')
2 df_test = pd.read_csv('test.csv')
```

▼ [3] Explore the Data

```
1 print(f'Training Data Shape Is: {df_train.shape}')
2 print(f'Test Data Shape Is: {df_test.shape}'')
```

→ Training Data Shape Is: (1460, 81)
Test Data Shape Is: (1459, 80)

```
1 print(f'The First 5 Rows In Train File:\n\n{df_train.head()}\n')
2 print('-' * 50)
3 print(f'\nThe First 5 Rows In Test File:\n\n{df_test.head()}\n')
```

→ The First 5 Rows In Train File:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	Norm
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	FR2	Gtl	Veenker	Feedr
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	CollgCr	Norm
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	Corner	Gtl	Crawfor	Norm
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	FR2	Gtl	NoRidge	Norm

The First 5 Rows In Test File:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	NAmes	Fee
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	AllPub	Corner	Gtl	NAmes	Nc

2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	AllPub	Inside	Gt1	Gilbert	Nc
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl	AllPub	Inside	Gt1	Gilbert	Nc
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS	AllPub	Inside	Gt1	StoneBr	Nc

```

1 print(f'\nThe Info Of Train File:\n')
2 df_train.info();
3 print('-' * 50)
4 print(f'\nThe Info Of Test File:\n')
5 df_test.info();

→ 24 Exterior2nd    1458 non-null   object
25 MasVnrType      565 non-null   object
26 MasVnrArea       1444 non-null   float64
27 ExterQual        1459 non-null   object
28 ExterCond        1459 non-null   object
29 Foundation       1459 non-null   object
30 BsmtQual         1415 non-null   object
31 BsmtCond         1414 non-null   object
32 BsmtExposure     1415 non-null   object
33 BsmtFinType1    1417 non-null   object
34 BsmtFinSF1       1458 non-null   float64
35 BsmtFinType2    1417 non-null   object
36 BsmtFinSF2       1458 non-null   float64
37 BsmtUnfSF        1458 non-null   float64
38 TotalBsmtSF      1458 non-null   float64
39 Heating           1459 non-null   object
40 HeatingQC         1459 non-null   object
41 CentralAir        1459 non-null   object
42 Electrical        1459 non-null   object
43 1stFlrSF          1459 non-null   int64
44 2ndFlrSF          1459 non-null   int64
45 LowQualFinSF     1459 non-null   int64
46 GrLivArea         1459 non-null   int64
47 BsmtFullBath     1457 non-null   float64
48 BsmtHalfBath     1457 non-null   float64
49 FullBath          1459 non-null   int64
50 HalfBath          1459 non-null   int64
51 BedroomAbvGr     1459 non-null   int64
52 KitchenAbvGr     1459 non-null   int64
53 KitchenQual       1458 non-null   object
54 TotRmsAbvGrd     1459 non-null   int64
55 Functional        1457 non-null   object
56 Fireplaces         1459 non-null   int64
57 FireplaceQu       729 non-null   object
58 GarageType         1383 non-null   object
59 GarageYrBlt       1381 non-null   float64
60 GarageFinish      1381 non-null   object
61 GarageCars         1458 non-null   float64
62 GarageArea         1458 non-null   float64
63 GarageQual        1381 non-null   object
64 GarageCond        1381 non-null   object
65 PavedDrive        1459 non-null   object
66 WoodDeckSF        1459 non-null   int64
67 OpenPorchSF       1459 non-null   int64
68 EnclosedPorch     1459 non-null   int64
69 3SsnPorch          1459 non-null   int64
70 ScreenPorch        1459 non-null   int64
71 PoolArea           1459 non-null   int64
72 PoolQC             3 non-null   object
73 Fence              290 non-null   object
74 MiscFeature        51 non-null   object
75 MiscVal            1459 non-null   int64
76 MoSold             1459 non-null   int64
77 YrSold             1459 non-null   int64
78 SaleType           1458 non-null   object
79 SaleCondition      1459 non-null   object
dtypes: float64(11), int64(26), object(43)
memory usage: 912.0+ KB

```

```

1 null_train_value = df_train.isnull().sum()
2 null_train_value = null_train_value[null_train_value > 0].sort_values(ascending=False)
3
4 null_test_value = df_test.isnull().sum()
5 null_test_value = null_test_value[null_test_value > 0].sort_values(ascending=False)
6
7 print(f'Null Values In Train File:\n\n{n{null_train_value}\n')
8 print('-' * 50)
9 print(f'\nNull Values In Test File:\n\n{n{null_test_value}\n')

```

```
→ Alley      1369
  Fence     1179
  MasVnrType 872
  FireplaceQu 690
  LotFrontage 259
  GarageType  81
  GarageYrBlt 81
  GarageFinish 81
  GarageQual   81
  GarageCond   81
  BsmtExposure 38
  BsmtFinType2 38
  BsmtQual     37
  BsmtCond     37
  BsmtFinType1 37
  MasVnrArea    8
  Electrical    1
dtype: int64
```

Null Values In Test File:

```
PoolQC      1456
MiscFeature 1408
Alley       1352
Fence       1169
MasVnrType  894
FireplaceQu 730
LotFrontage  227
GarageQual   78
GarageCond   78
GarageYrBlt  78
GarageFinish 78
GarageType   76
BsmtCond    45
BsmtQual    44
BsmtExposure 44
BsmtFinType1 42
BsmtFinType2 42
MasVnrArea   15
MSZoning     4
Functional    2
BsmtFullBath 2
Utilities     2
BsmtHalfBath 2
Exterior1st  1
Exterior2nd  1
TotalBsmtSF  1
BsmtUnfSF   1
BsmtFinSF2  1
BsmtFinSF1  1
KitchenQual  1
GarageArea    1
GarageCars    1
SaleType      1
dtype: int64
```

```
1 null_train_value_pct = (df_train.isnull().sum() / len(df_train)) * 100
2 null_train_value_pct = null_train_value_pct[null_train_value_pct > 0].sort_values(ascending=False)
3
4 null_test_value_pct = (df_test.isnull().sum() / len(df_test)) * 100
5 null_test_value_pct = null_test_value_pct[null_test_value_pct > 0].sort_values(ascending=False)
6
7 print(f'Null Values In Train File: % in training data:\n{null_train_value_pct}')
8 print('-' * 50)
9 print(f'Null Values In Test File: % in testing data:\n{null_test_value_pct}')
```

→ Null Values In Train File: % in training data:

```
PoolQC      99.520548
MiscFeature 96.301370
Alley       93.767123
Fence       80.753425
MasVnrType  59.726027
FireplaceQu 47.260274
LotFrontage 17.739726
GarageType  5.547945
GarageYrBlt 5.547945
GarageFinish 5.547945
GarageQual   5.547945
GarageCond   5.547945
```

```
BsmtExposure    2.602740
BsmtFinType2   2.602740
BsmtQual       2.534247
BsmtCond       2.534247
BsmtFinType1   2.534247
MasVnrArea     0.547945
Electrical      0.068493
dtype: float64
-----
Null Values In Test File: % in testing data:
PoolQC        99.794380
MiscFeature    96.504455
Alley          92.666210
Fence          80.123372
MasVnrType     61.274846
FireplaceQu    50.034270
LotFrontage    15.558602
GarageQual     5.346127
GarageCond     5.346127
GarageYrBlt    5.346127
GarageFinish    5.346127
GarageType      5.209047
BsmtCond       3.084304
BsmtQual       3.015764
BsmtExposure   3.015764
BsmtFinType1   2.878684
BsmtFinType2   2.878684
MasVnrArea     1.028101
MSZoning        0.274160
Functional      0.137080
BsmtFullBath   0.137080
Utilities        0.137080
BsmtHalfBath   0.137080
Exterior1st    0.068540
Exterior2nd    0.068540
TotalBsmtSF    0.068540
BsmtUnfSF      0.068540
BsmtFinSF2     0.068540
BsmtFinSF1     0.068540
KitchenQual    0.068540
GarageArea      0.068540
GarageCars      0.068540
SaleType         0.068540
dtype: float64
```

⌄ [4] Data Preprocessing & Cleaning

```
1 cols_to_drop = ['PoolQC', 'MiscFeature', 'Alley', 'Fence']
2
3 df_train = df_train.drop(columns=cols_to_drop)
4 df_test = df_test.drop(columns=cols_to_drop)
5
6 print("Training Data Shape After Dropping Columns:", df_train.shape)
7 print('-' * 50)
8 print("Test Data Shape After Dropping Columns:", df_test.shape)
```

⌄ Training Data Shape After Dropping Columns: (1460, 77)

```
-----
```

```
Test Data Shape After Dropping Columns: (1459, 76)
```

```
1 df_train['LotFrontage'] = df_train.groupby('Neighborhood')['LotFrontage'].transform(lambda x: x.fillna(x.median()))
2 df_test['LotFrontage'] = df_test.groupby('Neighborhood')['LotFrontage'].transform(lambda x: x.fillna(x.median()))
3
4 df_train['LotFrontage'] = df_train['LotFrontage'].fillna(df_train['LotFrontage'].median())
5 df_test['LotFrontage'] = df_test['LotFrontage'].fillna(df_test['LotFrontage'].median())
```

```
1 cols_to_fill_none = [
2     'GarageType', 'GarageFinish', 'GarageQual', 'GarageCond',
3     'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
4     'MasVnrType',
5     'FireplaceQu'
6 ]
7
8 df_train[cols_to_fill_none] = df_train[cols_to_fill_none].fillna('None')
9 df_test[cols_to_fill_none] = df_test[cols_to_fill_none].fillna('None')
```

```

1 print(df_train[cols_to_fill_none].isnull().sum())
2 print('-' * 50)
3 print(df_test[cols_to_fill_none].isnull().sum())

→ GarageType      0
  GarageFinish    0
  GarageQual     0
  GarageCond     0
  BsmtQual       0
  BsmtCond       0
  BsmtExposure   0
  BsmtFinType1   0
  BsmtFinType2   0
  MasVnrType     0
  FireplaceQu    0
  dtype: int64
-----
GarageType      0
GarageFinish    0
GarageQual     0
GarageCond     0
BsmtQual       0
BsmtCond       0
BsmtExposure   0
BsmtFinType1   0
BsmtFinType2   0
MasVnrType     0
FireplaceQu    0
dtype: int64

1 cat_cols = df_train.select_dtypes(include=['object']).columns
2
3 null_train_value_cat = df_train[cat_cols].isnull().sum()
4 null_test_value_cat = df_test[cat_cols].isnull().sum()
5
6 print("Remaining missing values in categorical columns (train):")
7 print(null_train_value_cat=null_train_value_cat > 0])
8 print("\nRemaining missing values in categorical columns (test):")
9 print(null_test_value_cat=null_test_value_cat > 0])

→ Remaining missing values in categorical columns (train):
  Electrical    1
  dtype: int64

Remaining missing values in categorical columns (test):
MSZoning        4
Utilities        2
Exterior1st     1
Exterior2nd     1
KitchenQual     1
Functional       2
SaleType         1
dtype: int64

1 another_cat_cols = [
2     'Electrical', 'MSZoning', 'Utilities', 'Exterior1st', 'Exterior2nd',
3     'KitchenQual', 'Functional', 'SaleType'
4 ]
5
6 for col in another_cat_cols:
7     if col in df_train.columns:
8         mode_value = df_train[col].mode()
9         mode_value = mode_value[0] if not mode_value.empty else "Unknown"
10
11     df_train[col] = df_train[col].fillna(mode_value)
12     df_test[col] = df_test[col].fillna(mode_value)
13
14 print(df_train[another_cat_cols].isnull().sum())
15 print('-' * 50)
16 print(df_test[another_cat_cols].isnull().sum())

→ Electrical      0
  MSZoning        0
  Utilities        0
  Exterior1st     0
  Exterior2nd     0
  KitchenQual     0
  Functional       0
  SaleType         0
  dtype: int64

```

```
-----
Electrical      0
MSZoning       0
Utilities       0
Exterior1st    0
Exterior2nd    0
KitchenQual    0
Functional      0
SaleType        0
dtype: int64
```

```
1 num_cols = [
2     'MasVnrArea',
3     'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
4     'BsmtFullBath', 'BsmtHalfBath',
5     'GarageCars', 'GarageArea',
6     'GarageYrBlt'
7 ]
8
9 for col in num_cols:
10    if col in df_train.columns:
11        df_train[col] = df_train[col].fillna(0)
12        df_test[col] = df_test[col].fillna(0)
13
14 print(df_train[num_cols].isnull().sum())
15 print('-' * 50)
16 print(df_test[num_cols].isnull().sum())
```

→ MasVnrArea 0
 BsmtFinSF1 0
 BsmtFinSF2 0
 BsmtUnfSF 0
 TotalBsmtSF 0
 BsmtFullBath 0
 BsmtHalfBath 0
 GarageCars 0
 GarageArea 0
 GarageYrBlt 0
 dtype: int64

```
-----
MasVnrArea      0
BsmtFinSF1      0
BsmtFinSF2      0
BsmtUnfSF       0
TotalBsmtSF     0
BsmtFullBath    0
BsmtHalfBath    0
GarageCars       0
GarageArea       0
GarageYrBlt     0
dtype: int64
```

```
1 print(f'Any Null values in df_train:\n{df_train.isnull().sum().sum()}')
2 print('-' * 50)
3 print(f'Missing values in df_test:\n{df_test.isnull().sum().sum()}')
```

→ Any Null values in df_train:

0

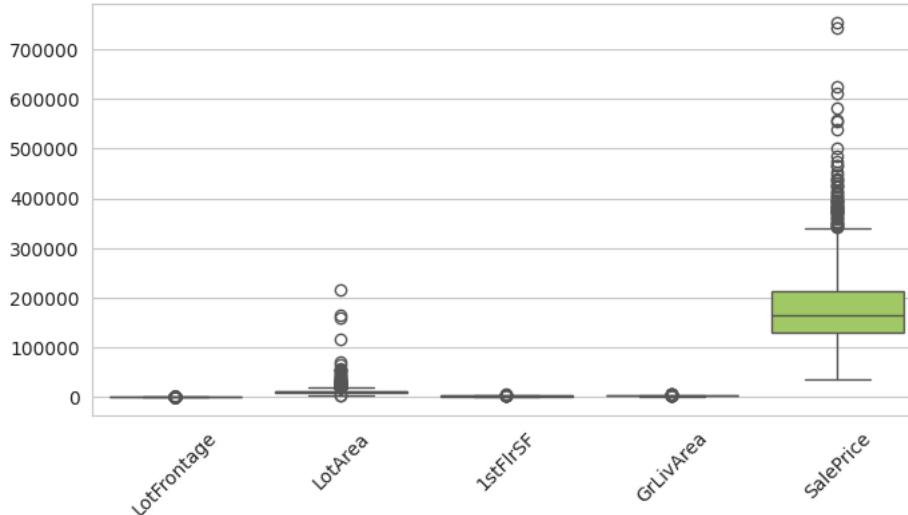
Missing values in df_test:

0

```
1 num_features = ['LotFrontage', 'LotArea', '1stFlrSF', 'GrLivArea', 'SalePrice']
2
3 plt.figure(figsize=(8, 4))
4 sns.boxplot(data=df_train[num_features])
5 plt.title('Boxplots for Numerical Features (Outlier Detection)')
6 plt.xticks(rotation=45)
7 plt.show()
```



Boxplots for Numerical Features (Outlier Detection)



```

1 skewed_features = ['LotArea', 'GrLivArea', '1stFlrSF', 'TotalBsmtSF']
2
3 for feature in skewed_features:
4     df_train[feature] = np.log1p(df_train[feature])
5     df_test[feature] = np.log1p(df_test[feature])
6
7 df_train['SalePrice'] = np.log1p(df_train['SalePrice'])

1 x = df_train.drop(columns=['Id', 'SalePrice'])
2 y = df_train['SalePrice']
3
4 x_train, x_val, y_train, y_val = train_test_split(x, y, test_size=0.2, random_state=2025)

1 categorical_cols = x_train.select_dtypes(include=['object']).columns.tolist()
2 numerical_cols = x_train.select_dtypes(include=['int64', 'float64']).columns.tolist()
3
4 preprocessor_standard = ColumnTransformer([
5     ('num', StandardScaler(), numerical_cols),
6     ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_cols)
7 ])
8
9 preprocessor_minmax = ColumnTransformer([
10    ('num', MinMaxScaler(), numerical_cols),
11    ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_cols)
12 ])

```

▼ [05] Build & Evaluate Models Using Pipeline

```

1 models = [
2     ('Dummy', DummyRegressor()),
3     ('Linear', LinearRegression()),
4     ('Ridge', Ridge()),
5     ('Lasso', Lasso()),
6     ('Poly', Pipeline([
7         ('poly', PolynomialFeatures(degree=2, include_bias=False)),
8         ('linear', LinearRegression())
9     ])),
10    ('DecisionTree', DecisionTreeRegressor(random_state=2025)),
11    ('RandomForest', RandomForestRegressor(random_state=2025)),
12    ('XGBoost', XGBRegressor(random_state=2025, verbosity=0)),
13    ('LightGBM', LGBMRegressor(random_state=2025)),
14    ('SVR', SVR()),
15    ('KNN', KNeighborsRegressor()),
16    ('NeuralNet', MLPRegressor(max_iter=500, random_state=2025))
17 ]

```

```

1 def evaluate_model(model, preprocessor, X_train, y_train, X_val, y_val):
2     pipeline = Pipeline([

```

```

3     ('preprocessor', preprocessor),
4     ('model', model)
5 )
6
7 pipeline.fit(X_train, y_train)
8 y_pred = pipeline.predict(X_val)
9
10 mae = mean_absolute_error(y_val, y_pred)
11 rmse = np.sqrt(mean_squared_error(y_val, y_pred))
12 r2 = r2_score(y_val, y_pred)
13
14 return mae, rmse, r2

1 results = []
2
3 for name, model in models:
4     for preprocessor_name, preprocessor in [
5         ('StandardScaler', preprocessor_standard),
6         ('MinMaxScaler', preprocessor_minmax)
7     ]:
8         mae, rmse, r2 = evaluate_model(model, preprocessor, x_train, y_train, x_val, y_val)
9
10    results.append({
11        'Model': name,
12        'Scaler': preprocessor_name,
13        'MAE': mae,
14        'RMSE': rmse,
15        'R2': r2
16    })
17
18 print(f"Model: {name:<12} | Scaler: {preprocessor_name:<15} | MAE: {mae:.4f} | RMSE: {rmse:.4f} | R2: {r2:.4f}")

Model: Dummy | Scaler: StandardScaler | MAE: 0.3020 | RMSE: 0.3982 | R2: -0.0006
Model: Dummy | Scaler: MinMaxScaler | MAE: 0.3020 | RMSE: 0.3982 | R2: -0.0006
Model: Linear | Scaler: StandardScaler | MAE: 0.0938 | RMSE: 0.1850 | R2: 0.7841
Model: Linear | Scaler: MinMaxScaler | MAE: 0.0938 | RMSE: 0.1850 | R2: 0.7841
Model: Ridge | Scaler: StandardScaler | MAE: 0.0907 | RMSE: 0.1761 | R2: 0.8042
Model: Ridge | Scaler: MinMaxScaler | MAE: 0.0909 | RMSE: 0.1762 | R2: 0.8040
Model: Lasso | Scaler: StandardScaler | MAE: 0.3020 | RMSE: 0.3982 | R2: -0.0006
Model: Lasso | Scaler: MinMaxScaler | MAE: 0.3020 | RMSE: 0.3982 | R2: -0.0006
Model: Poly | Scaler: StandardScaler | MAE: 0.1157 | RMSE: 0.1646 | R2: 0.8291
Model: Poly | Scaler: MinMaxScaler | MAE: 0.1218 | RMSE: 0.1888 | R2: 0.7749
Model: DecisionTree | Scaler: StandardScaler | MAE: 0.1480 | RMSE: 0.2026 | R2: 0.7409
Model: DecisionTree | Scaler: MinMaxScaler | MAE: 0.1496 | RMSE: 0.2025 | R2: 0.7412
Model: RandomForest | Scaler: StandardScaler | MAE: 0.0956 | RMSE: 0.1539 | R2: 0.8505
Model: RandomForest | Scaler: MinMaxScaler | MAE: 0.0956 | RMSE: 0.1539 | R2: 0.8505
Model: XGBoost | Scaler: StandardScaler | MAE: 0.0960 | RMSE: 0.1488 | R2: 0.8602
Model: XGBoost | Scaler: MinMaxScaler | MAE: 0.0984 | RMSE: 0.1528 | R2: 0.8526
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001148 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3229
[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 186
[LightGBM] [Info] Start training from score 12.026044
Model: LightGBM | Scaler: StandardScaler | MAE: 0.0882 | RMSE: 0.1439 | R2: 0.8693
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.001286 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 3205
[LightGBM] [Info] Number of data points in the train set: 1168, number of used features: 186
[LightGBM] [Info] Start training from score 12.026044
Model: LightGBM | Scaler: MinMaxScaler | MAE: 0.0880 | RMSE: 0.1429 | R2: 0.8711
Model: SVR | Scaler: StandardScaler | MAE: 0.0967 | RMSE: 0.1631 | R2: 0.8321
Model: SVR | Scaler: MinMaxScaler | MAE: 0.0974 | RMSE: 0.1555 | R2: 0.8473
Model: KNN | Scaler: StandardScaler | MAE: 0.1271 | RMSE: 0.1986 | R2: 0.7510
Model: KNN | Scaler: MinMaxScaler | MAE: 0.1553 | RMSE: 0.2166 | R2: 0.7040
Model: NeuralNet | Scaler: StandardScaler | MAE: 0.1298 | RMSE: 0.2451 | R2: 0.6208
Model: NeuralNet | Scaler: MinMaxScaler | MAE: 0.1421 | RMSE: 0.2411 | R2: 0.6330
```

✓ [6] Improving Models with GridSearchCV

The best 3 models have been selected:

1. LightGBM (MinMaxScaler) ==> [MAE: 0.0880 | RMSE: 0.1429 | R2: 0.8711]
2. XGBoost (StandardScaler) ==> [MAE: 0.0960 | RMSE: 0.1488 | R2: 0.8602]
3. RandomForest (StandardScaler) ==> [MAE: 0.0956 | RMSE: 0.1539 | R2: 0.8505]

These models will be optimized by **GridSearchCV** to select the best ones.

```
1 param_grid_lgb = {
2     'model__n_estimators': [100, 200],
3     'model__learning_rate': [0.05, 0.1],
4     'model__max_depth': [5, 7],
5     'model__num_leaves': [31, 63]
6 }
7
8 pipeline_lgb = Pipeline([
9     ('preprocessor', preprocessor_minmax),
10    ('model', LGBMRegressor(random_state=2025))
11 ])
12
13 grid_lgb = GridSearchCV(
14     pipeline_lgb,
15     param_grid_lgb,
16     cv=5,
17     scoring='neg_mean_squared_error',
18     n_jobs=-1,
19     verbose=1
20 )
21
22 grid_lgb.fit(x_train, y_train)
23
24 print("Best transactions:", grid_lgb.best_params_)
25 print("Best RMSE on verification:", round(np.sqrt(-grid_lgb.best_score_), 4))
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

1 param_grid_xgb = {
2     'model__n_estimators': [100, 200],
3     'model__learning_rate': [0.05, 0.1],
4     'model__max_depth': [5, 7],
5     'model__subsample': [0.8, 1.0]
6 }
7
8 pipeline_xgb = Pipeline([
9     ('preprocessor', preprocessor_standard),
10    ('model', XGBRegressor(random_state=2025, verbosity=0))
11 ])
12
13 grid_xgb = GridSearchCV(
14     pipeline_xgb,
15     param_grid_xgb,
16     cv=5,
17     scoring='neg_mean_squared_error',
18     n_jobs=-1,
19     verbose=1
20 )
21
22 grid_xgb.fit(x_train, y_train)
23
24 print("Best transactions:", grid_xgb.best_params_)
25 print("Best RMSE on verification:", round(np.sqrt(-grid_xgb.best_score_), 4))
```

→ Fitting 5 folds for each of 16 candidates, totalling 80 fits
 Best transactions: {'model__learning_rate': 0.1, 'model__max_depth': 5, 'model__n_estimators': 200, 'model__subsample': 0.8}
 Best RMSE on verification: 0.1315

```
1 param_grid_rf = {
2     'model__n_estimators': [100, 200],
3     'model__max_depth': [10, 20, None],
4     'model__min_samples_split': [2, 5],
5     'model__min_samples_leaf': [1, 2]
6 }
7
8 pipeline_rf = Pipeline([
9     ('preprocessor', preprocessor_standard),
10    ('model', RandomForestRegressor(random_state=2025))
11 ])
12
13 grid_rf = GridSearchCV(
14     pipeline_rf,
15     param_grid_rf,
16     cv=5,
17     scoring='neg_mean_squared_error',
18     n_jobs=-1,
19     verbose=1
20 )
21
22 grid_rf.fit(x_train, y_train)
23
24 print("Best transactions:", grid_rf.best_params_)
25 print("Best RMSE on verification:", round(np.sqrt(-grid_rf.best_score_), 4))
```

→ Fitting 5 folds for each of 24 candidates, totalling 120 fits
 Best transactions: {'model__max_depth': None, 'model__min_samples_leaf': 2, 'model__min_samples_split': 2, 'model__n_estimators': 200}
 Best RMSE on verification: 0.1424

✓ [7] Prediction on test file "df_test"

- The "XGBoost" model was selected as the best model based on previous optimization results.
- The selected model will be used to predict unit prices based on the data in the **test file**.
- The results will be printed in an Excel file titled "**Submission**".

```
1 test_ids = df_test['Id']
2
3 x_test = df_test.drop(columns=['Id'])
```

```

4
5 y_test_pred_log = grid_xgb.predict(x_test)
6
7 y_test_pred = np.expm1(y_test_pred_log)
8
9 submission = pd.DataFrame({
10     'Id': test_ids,
11     'SalePrice': y_test_pred.round(2)
12 })
13
14 submission.to_csv('submission.csv', index=False)
15
16 print("The submission.csv file was created successfully.")
17 print(submission.head())

```

→ The submission.csv file was created successfully.

		SalePrice
0	1461	128137.757812
1	1462	156616.250000
2	1463	185175.296875
3	1464	186597.406250
4	1465	179302.843750

```

1 submission.to_excel('submission.xlsx', index=False)
2
3 files.download('submission.xlsx')

```

→

▼ [8] Building a Streamlit Model

```

1 joblib.dump(grid_xgb, 'house_price_model.pkl')
2 print("The model has been saved.")
3
4 model = joblib.load('house_price_model.pkl')
5
6 def predict_price(total_sqft, bedrooms, bathrooms, garage_cars, overall_qual, year_built, neighborhood):
7     input_data = pd.DataFrame({
8         'MSSubClass': [60],
9         'MSZoning': ['RL'],
10        'LotFrontage': [80],
11        'LotArea': [total_sqft * 2],
12        'Street': ['Pave'],
13        'Alley': ['None'],
14        'LotShape': ['Reg'],
15        'LandContour': ['Lvl'],
16        'Utilities': ['AllPub'],
17        'LotConfig': ['Inside'],
18        'LandSlope': ['Gtl'],
19        'Neighborhood': [neighborhood],
20        'Condition1': ['Norm'],
21        'Condition2': ['Norm'],
22        'BldgType': ['1Fam'],
23        'HouseStyle': ['2Story'],
24        'OverallQual': [overall_qual],
25        'OverallCond': [5],
26        'YearBuilt': [year_built],
27        'YearRemodAdd': [year_built],
28        'RoofStyle': ['Gable'],
29        'RoofMatl': ['CompShg'],
30        'Exterior1st': ['VinylSd'],
31        'Exterior2nd': ['VinylSd'],
32        'MasVnrType': ['None'],
33        'MasVnrArea': [0],
34        'ExterQual': ['TA'],
35        'ExterCond': ['TA'],
36        'Foundation': ['PConc'],
37        'BsmtQual': ['TA'],
38        'BsmtCond': ['TA'],
39        'BsmtExposure': ['No'],
40        'BsmtFinType1': ['GLQ'],
41        'BsmtFinSF1': [total_sqft * 0.3],
42        'BsmtFinType2': ['Unf'],
43        'BsmtFinSF2': [0],
44        'BsmtUnfSF': [total_sqft * 0.2],

```

```

45     'TotalBsmtSF': [total_sqft * 0.5],
46     'Heating': ['GasA'],
47     'HeatingQC': ['Ex'],
48     'CentralAir': ['Y'],
49     'Electrical': ['SBrkr'],
50     '1stFlrSF': [total_sqft * 0.6],
51     '2ndFlrSF': [total_sqft * 0.4],
52     'LowQualFinSF': [0],
53     'GrLivArea': [total_sqft],
54     'BsmtFullBath': [1],
55     'BsmtHalfBath': [0],
56     'FullBath': [bathrooms],
57     'HalfBath': [1 if bathrooms >= 2 else 0],
58     'BedroomAbvGr': [bedrooms],
59     'KitchenAbvGr': [1],
60     'KitchenQual': ['TA'],
61     'TotRmsAbvGrd': [bedrooms + bathrooms],
62     'Functional': ['Typ'],
63     'Fireplaces': [1],
64     'FireplaceQu': ['TA'],
65     'GarageType': ['Attchd'],
66     'GarageYrBlt': [year_built],
67     'GarageFinish': ['RFn'],
68     'GarageCars': [garage_cars],
69     'GarageArea': [garage_cars * 300],
70     'GarageQual': ['TA'],
71     'GarageCond': ['TA'],
72     'PavedDrive': ['Y'],
73     'WoodDeckSF': [200],
74     'OpenPorchSF': [50],
75     'EnclosedPorch': [0],
76     '3SsnPorch': [0],
77     'ScreenPorch': [30],
78     'PoolArea': [0],
79     'PoolQC': ['None'],
80     'Fence': ['None'],
81     'MiscFeature': ['None'],
82     'MiscVal': [0],
83     'MoSold': [6],
84     'YrSold': [2024],
85     'SaleType': ['WD'],
86     'SaleCondition': ['Normal']
87   })
88
89 try:
90   prediction_log = model.predict(input_data)
91   prediction = np.expm1(prediction_log)
92   return f"${prediction[0]:,.2f}"
93 except Exception as e:
94   return f"Mistake: {e}"
95
96 interface = gr.Interface(
97   fn=predict_price,
98   inputs=[
99     gr.Number(label="Total area (square feet"),
100    gr.Slider(1, 10, step=1, label="Number of rooms"),
101    gr.Slider(1, 6, step=1, label="Number of bathrooms"),
102    gr.Slider(0, 4, step=1, label="Garage capacity (number of cars"),
103    gr.Slider(1, 10, step=1, label="Build Quality (1 to 10)"),
104    gr.Number(1800, 2024, step=1, label="Year of construction"),
105    gr.Dropdown(['CollgCr', 'Veenker', 'Crawfor', 'NoRidge', 'Mitchel', 'Somerst', 'OldTown', 'BrkSide', 'NAmes'], label="The neighborhood"),
106  ],
107  outputs=gr.Textbox(label="Expected price"),
108  title="Home price forecast",
109  description="Enter the home's specifications to get the estimated price."
110 )
111
112 interface.launch()

```

The model has been saved.

It looks like you are running Gradio on a hosted Jupyter notebook, which requires `share=True`. Automatically setting `share=True` (you

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

* Running on public URL: <https://1751d2e3db5ff482da.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working dir

Home price forecast

Enter the home's specifications to get the estimated price.

Total area (square feet)	100	Expected price	\$341,045.69
Number of rooms	4	Flag	
1	10		
Number of bathrooms	2		
1	6		
Garage capacity (number of cars)	3		
0	4		
Build Quality (1 to 10)	8		
1	10		
..	..		